

Programming Paradigms

Logic Programming using Prolog

Department of Computer Science and Information Systems
Birkbeck, University of London

October 8, 2014



Logic Programming

- Programming languages for logic programming are very different to those encountered so far — they are *declarative* languages
- You present the facts and inference rules and the program will do the reasoning
- In a declarative language
 - the programmer specifies a goal to be achieved
 - the system then *works out* how to achieve it
- In imperative and object-oriented languages, the programmer has to do both

Prolog


One of the most well-known logic programming languages is Prolog

- Stands for **P**rogrammation en **L**ogique (Programming in Logic)
- Developed by Alain Colmerauer and colleagues in the early 1970s
- University of Edinburgh a major player (Clocksin and Mellish) together with Imperial College

Relational databases owe something to Prolog




Alain Colmerauer



Robert Kowalski

The Book



Mathematical Foundation

Just a brief explanation how Prolog fits into the framework of mathematical logic

- *First-order logic* is a powerful mathematical tool for formalizing descriptions
 - It is also sometimes called *predicate logic*
- Unfortunately, first-order logic is not **decidable**
- Prolog is based on a decidable subset of first-order logic called *Horn clauses*
- It is still Turing-complete, though

First Program

Let's have a look at a very simple program: Hello World!

```
?- write('Hello World!'), nl.
```

with output¹:

```
Hello World!
```

```
yes
```

Although this works, it's an atypical example of a Prolog program

¹(may say true instead of yes dependent upon the version of Prolog used)

Language Basics

Prolog has two aspects:

- One to express the data
- Another to query the data

Data is represented in the form of facts and logical rules

Facts: a fact is a basic assertion about some world
e.g., Babe is a pig; pigs like mud

Rules: a rule is an inference about facts in that world
e.g., An animal likes mud if it is a pig

Query: a query is a question about that world
e.g., Does Babe like mud?

Facts and rules go into a *knowledge base*

- Prolog allows you to express the contents of a knowledge base
- Usually a compiler turns this base into a form efficient for querying

Querying links together facts and rules to tell you something about the world modeled in the knowledge base

Simple Knowledge Base and Queries I

First some remarks about syntax:

- If a word begins with a lower-case character, it's an *atom*
- An atom is a fixed value, similar to a Ruby symbol
- If it begins with an upper-case letter (or an underscore), it's a *variable*

Simple Knowledge Base and Queries II



Let's have a look at a very simple knowledge base

```
likes(wallace, toast).  
likes(wallace, cheese).  
likes(gromit, cheese).  
likes(gromit, cake).  
likes(wendolene, sheep).  
friend(X,Y) :- likes(X,Z),likes(Y,Z),\+(X=Y).
```

The first five statements are facts, the last one is a rule

Facts I

- In the facts on the previous slide, wallace, gromit, wendolene, toast, cheese, cake, and sheep are atoms
- The facts can be read as

“Wallace likes toast”

“Wallace likes cheese”

“Gromit likes cheese”

“Gromit likes cake”

“Wendolene likes sheep”

Facts II

- The name of a relationship (before the round brackets) is called a *predicate* (e.g., the predicate `likes` has two parameters)
- The order of atoms in a fact is important, e.g., “cheese likes Wallace” is not a fact

We are now ready to ask some questions. . .

Queries I

- The most basic queries are questions about facts with a *yes/no* answer
- The following queries are quite intuitive where Prolog tries to match a query to known facts

```
?- likes(wallace,sheep).  
no  
?- likes(gromit,cheese).  
yes
```

- Not very exciting — Prolog is just throwing the facts back at us

Let's try something else...

Queries II

Some atom that isn't in the knowledge base...

```
?- likes(fluffles,sausage).  
no
```

So no actually means that Prolog cannot prove this statement given the current state of the knowledge base

Instantiation I

- We can ask Prolog to find values for variables:

```
?- likes(Who,cheese).
```

- Who is an *uninstantiated* variable, i.e., it does not have a value assigned to it
- Prolog searches the knowledge base from the beginning trying to find a matching fact
- The first matching fact found is `likes(wallace,cheese)`, so Who is *instantiated* with wallace
- At this point Prolog outputs `Who = wallace`, stops, and asks us what to do

Instantiation II

We can then either

- (i) stop searching by just hitting the return key, or
- (ii) continue searching by entering ;

If we continue, Prolog

- (i) forgets the value `wallace` for the variable `Who`
- (ii) and continues at the position it previously stopped

Continuing will output `Who = gromit` and then `no` (when it finds no further solutions)

Goals I

- By submitting a query, we ask Prolog to try to satisfy a *goal*
- We can ask Prolog to satisfy the conjunction of two goals:

```
?- likes(wallace,toast),likes(gromit,toast).  
no
```

- We can combine conjunctions with variables to make queries more interesting
- Now that we found out that at least one of them does not like toast...

Goals II

- ...is there something both of them like?

```
?- likes(wallace,What),likes(gromit,What).  
What = cheese ? ;  
no
```

- How does Prolog process this query (conceptually)?
- It uses **backtracking** to try to satisfy the first goal and then the second goal

Backtracking I

likes(wallace,What), likes(gromit,What).
first goal second goal

likes(wallace,cheese).


likes(wallace,toast).

likes(gromit,cheese).


likes(gromit,cake).

likes(wendolene,sheep).


Backtracking II




Backtracking III



Backtracking IV



Backtracking V



Rules I

- A *rule* is a general statement about objects and their relationships
- A rule in plain English could be:
X is a sister of Y if:
X is female and
X and Y have the same parents.
- Important: a variable stands for the same object wherever it occurs in a rule

Rules II

- Rules in Prolog consist of a *head* and a *body* connected by the symbol :- (which is pronounced *if*)
- The predicate `friend` in our knowledge base is defined by a rule
- Predicates can be defined by a combination of facts and rules
- A clause of a predicate is a fact or rule defining the predicate

Rules III

- If we want to express that

“Wallace is a friend of anyone who likes cheese”

we could formulate it like this:

```
friend(wallace,X) :- likes(X,cheese).
```

- Running the query `friend(wallace,X).` will produce two results: `wallace` and `gromit`
- We can exclude `wallace` by saying that `X` shouldn't be `wallace`

Rules IV

```
friend(wallace,X) :- likes(X,cheese),  
                    \+(X=wallace).
```

- $\backslash +$ is the negation of a subgoal
- This only lists friends of wallace (those persons who like cheese)
- A generalization of the rule would be

“X and Y are friends, if X and Y like the same Z and X and Y are not the same”

```
friend(X,Y) :- likes(X,Z),likes(Y,Z),\+(X=Y).
```

Rules V

Let's try it out:

```
?- friend(gromit,wallace).  
yes  
?- friend(wallace,gromit).  
yes  
?- friend(wallace,wallace).  
no  
?- friend(wallace,wendolene).  
no
```

Rules VI

Now let's ask who is a friend of Wallace:

```
?- friend(wallace,Who).  
Who = gromit ? ;  
no
```

Or let's find all pairs of friends:

```
?- friend(Who1,Who2).  
  
Who1 = wallace  
Who2 = gromit ? ;  
  
Who1 = gromit  
Who2 = wallace ? ;  
  
no
```

Another example: from Bratko

...

Example — The Map Colouring Problem I

Just using facts, rules, and variables we can already do some interesting things



Assume we want to colour a map, such that two regions with a common border don't have the same colour

Example — The Map Colouring Problem II

- In order to simplify things, we'll only look at regions 3,4,5, and 6 and use the colours red, green, and blue
- Now all we have to do is describe this to Prolog:

```
border(r,g) . border(r,b) .  
border(g,r) . border(g,b) .  
border(b,r) . border(b,g) .
```

```
colouring(L,TAA,V,FVG) :-  
    border(L,TAA) ,  
    border(L,V) ,  
    border(TAA,V) ,  
    border(V,FVG) .
```


Example — The Map Colouring Problem III

Querying `colouring(L,TAA,V,FVG)` . will now provide all the answers:

```
?- colouring(L,TAA,V,FVG) .
```

```
FVG = r
```

```
L = r
```

```
TAA = g
```

```
V = b ? ;
```

```
FVG = g
```

```
L = r
```

```
TAA = g
```

```
V = b ?
```

```
...
```

Where's the Program?

In Prolog you don't have to write a program

- You express the logic of a problem in facts and inferences
- And then let the computer do the work in figuring out a solution

Solving the map colouring problem with a language like Java or Ruby would be much harder to do

Anonymous Variables

- Sometimes we want to use a variable but don't care with which value it is instantiated (We don't want to use the variable anywhere else)
- For example, we want to find out if there is anyone who likes cheese (but we don't need to know who)

```
?- likes(_,cheese).  
true ?  
yes
```

- We use an underscore `_` for the anonymous variable
- Several occurrence of `_` in the same clause do not need to be given consistent interpretations

Structures I

- If we want to say that Wallace and Wendolene own books, we could formulate the following facts

```
owns(wallace, book).  
owns(wendolene, book).
```

- However, this means that Wallace owns the same object that Wendolene owns
- Specifying the title to distinguish may not help:

```
owns(wallace, perfume).  
owns(wendolene, russell_the_sheep).
```

- It's not clear we are talking about books here

Structures II

- We can introduce a *structure* for books
- A structure is the closest construct to a data type in Prolog
- A term can be decomposed into its components:

```
owns(wallace,book(perfume,suesskind)).  
owns(wendolene,book(russell_the_sheep,scotton)).
```

- Looking at `book(perfume,suesskind)`
 - `book` is the *functor* of the structure
 - `perfume` and `suesskind` are its *components*

Structures III

- Structures can be nested (arbitrarily deep):

```
owns(gromit,book(wuthering_heights,  
                 author(emily,bronte))).
```

- We can use structures in querying:
- For example, if we want to know if Gromit owns any books written by one of the Brontë sisters, we would query:

```
?- owns(gromit,book(X,author(Y,bronte))).
```

```
X = wuthering_heights
```

```
Y = emily
```

Structures IV

- The syntax for structures looks identical to that for facts
- A predicate is actually the functor of a structure
- The arguments of a fact or rule are components of a structure

Equality and Matching I

- Prolog has a number of built-in predicates
- One of them is *equality* written as “=”
- Following Prolog syntax, it should be written as $=(X,Y)$
 - While the above works, Prolog also allows you to use an infix notation: $X=Y$
- Prolog attempts to match X and Y , the goal succeeds if they match

Equality and Matching II

- Integers and atoms are always equal to themselves:

```
?- wallace = wallace.  
yes  
?- cheese = cake.  
no  
?- 1066 = 1066.  
yes  
?- 1206 = 1583.  
no
```

- A variable always matches itself:

```
?- X = X.  
yes
```

Equality and Matching III

- If we match two different variables, e.g. $X = Y$, we have to distinguish two cases
 - 1 None or one variable is instantiated
 - 2 Both are instantiated

Case 1: as soon as one them is instantiated with a value, the other will be instantiated with the same value

```
?- X = Y, likes(X,toast).  
X = wallace  
Y = wallace ?  
yes
```

Equality and Matching IV

Case 2: if both are already instantiated, then it depends on the value they are instantiated with

```
?- likes(X,cheese),likes(Y,cake),X=Y.  
X = gromit  
Y = gromit ?  
yes  
?- likes(X,toast),likes(Y,cake),X=Y.  
no
```

Comparison and Matching

- Prolog also offers other comparison operators:

```
?- 2 > 3.  
no  
?- 3 >= 2.  
yes  
?- 3 =< 2.  
no  
?- X \= Y.  
no
```

- The last one means that X cannot be made equal to Y
- You cannot redefine built-in predicates, stating the following as a fact will raise an error

```
2 > 3.
```

Arithmetic I

- Prolog also offers the standard arithmetic operators: $+$, $-$, $*$, $/$, mod ,
- Just typing in an arithmetic operation will not actually carry it out

```
?- 7 = 3 + 4.  
no
```

- Using the `is` operator will evaluate the right-hand side and match it to the left-hand side

```
?- 7 is 3 + 4.  
yes
```

Arithmetic II

- Given the following fact base, compute the population density of countries:

```
pop(usa,313).  
pop(italy,61).  
pop(uk,63).  
area(usa,9.826).  
area(italy,0.301).  
area(uk,0.243).
```

- The following rule computes the density:

```
density(X,Y) :- pop(X,P),area(X,A),Y is P/A.
```

Arithmetic III

- Compute population density of USA:

```
?- density(usa,Y).  
Y = 31.854264197028289  
yes
```

- Compute all densities:

```
?- density(X,Y).  
X = usa  
Y = 31.854264197028289 ? ;  
X = italy  
Y = 202.65780730897012 ? ;  
X = uk  
Y = 259.25925925925924  
yes
```

Lists I

- We have already seen structures as a construct to build more complicated data types
- Another important type supported by Prolog is a list
- The elements of a list are enclosed in square brackets:

```
?- [1,2,3] = [1,2,3] .
```

```
yes
```

```
?- [1,2,3] = [X,Y,Z] .
```

```
X = 1
```

```
Y = 2
```

```
Z = 3
```

```
yes
```


Lists II

- We can split lists into a *head* and *tail* using the “|” operator:

```
?- [Head|Tail] = [1,2,3].
```

```
Head = 1
```

```
Tail = [2,3]
```

```
yes
```

```
?- [Head|Tail] = [].
```

```
no
```

```
?- [Head|Tail] = [1].
```

```
Head = 1
```

```
Tail = []
```

```
yes
```

Recursion

- Let's assume we want to find out if an element is part of a list
- We have to do this recursively in Prolog
- Recursion in Prolog means that a predicate appears on the left- and the right-hand side of a rule
- For example: an element is *in* a list if it is
 - the head of the list
 - *in* the tail of the list

```
is_in(X,[X|_]).  
is_in(X,[_|Y]) :- is_in(X,Y).  
  
?- is_in(d,[a,b,c,d,e,f]).  
true
```

Let us take a closer Look I

- You might have noticed that in the book by Bruce Tate, the friend rule was written differently:

```
friend(X, Y) :- \+(X=Y),likes(X,Z),likes(Y,Z).
```

- Might not look like a big change, but this has consequences
- For example, if we run the query `friend(wallace,Y).` with the above rule, we get

```
?- friend(wallace,Y).  
no
```

- What is going on here?

Let us take a closer Look II

- The position of the predicate $\backslash+(X=Y)$ has a big impact
- Prolog tries to satisfy subgoals from left to right
- $\backslash+(X=Y)$ fails if $X=Y$ can be satisfied
 - 1 X and Y start off uninstantiated in the above case
 - 2 As soon as one of them is instantiated, the other will take on the same value
 - 3 This makes $X=Y$ true, resulting in $\backslash+(X=Y)$ being false
 - 4 Consequently, the first subgoal always fails

Let us take a closer Look III

- If we arrange the predicates in a different order

```
friend(X, Y) :- likes(X,Z),likes(Y,Z),\+(X=Y).
```

- then X and Y will already be instantiated when reaching the subgoal $\backslash+(X=Y)$
- If X and Y have a different value at that point, then $\backslash+(X=Y)$ will succeed
- It is important to get the order right in which variables are instantiated!

“Cutting” the Number of Solutions I

If you ask Prolog to keep looking for further solutions (by answering with ;) it will go through all possible solutions using backtracking:

```
dance_pairs(X,Y) :- boy(X), girl(Y).  
boy(adam).  
boy(bert).  
...  
girl(angela).  
girl(betty).  
...  
?- dance_pairs(X,Y).  
X = adam, Y = angela ;  
X = adam, Y = betty ;  
...
```

“Cutting” the Number of Solutions II

Sometimes we are not interested in exhaustively going through all solutions:

- We only want to know if a solution exists
- We are happy with a certain subset of solutions
- In some recursive cases, there may be an infinite number of solutions

Prolog provides the **cut** operator to force it not to consider certain choices

The Cut Operator I

- The cut operator is denoted by ! and can be inserted into a rule as a subgoal
- What does it do? Let's have a look:

```
foo :- a,b.  
foo :- c,d,!,e,f.  
foo :- g,h.
```

- First of all, ! always succeeds, i.e., if c and d are satisfied in the second rule, then Prolog will immediately start matching e
- But there's more to it...

The Cut Operator II

- Assuming c and d are satisfied while checking the second rule, then the choices made for c and d are “locked in”
 - Prolog may not go back and search for other solutions for c and d
 - It may still do backtracking for e and f , though
- In addition to this, if the second rule fails, Prolog may not go beyond this rule to try to satisfy foo
 - It will not try out $foo \text{ :- } g, h$.
- How is the cut operator used in practice?

Confirming choice of a rule I

The first use is to tell Prolog that it has found the right rule to apply

- Assume we want to add up the numbers from 1 to N

```
sum_to(1,1).  
sum_to(N,Result) :- TmpN is N-1,  
                    sum_to(TmpN,TmpRes),  
                    Result is TmpRes + N.
```

- While this works, it may start an infinite recursion:

```
?- sum_to(3,X).  
X = 6 ? ;  
Fatal Error: local stack overflow
```

Confirming choice of a rule II

- Asking for another solution forces Prolog to search for another solution for `sum_to(1,TmpRes)`, applying the rule `sum_to(N,Result)` to it
- Applying the rule will search for a solution for `sum_to(0,TmpRes)`, which in turn will again apply the rule
- Next attempt at satisfying will be to try to match `sum_to(-1,TmpRes)` and so on

Confirming choice of a rule III

- We want to tell Prolog that once it has matched the fact `sum_to(1,1)`. it should not try searching for further solutions
- We can achieve this by rewriting the fact:

```
sum_to(1,1) :- !.  
sum_to(N,Result) :- TmpN is N-1,  
                    sum_to(TmpN,TmpRes),  
                    Result is TmpRes + N.
```

```
?- sum_to(3,X).  
X = 6  
yes
```

Confirming choice of a rule IV

- We could just tell Prolog to stop searching for further solutions in the above example
- However, this may not always be under our control

```
go :- sum_to(1,X), foo(apples).  
?- go.
```

- If `foo(apples)` fails, then this will trigger backtracking on `sum_to(1,X)`

“Cut-Fail” Combination I

The second use of the cut operator involves the built-in `fail` predicate that cannot be satisfied:

```
p(X) :- fail.  
?- p(X).  
no
```

Let us consider an example which tries to figure out the correct tax rate for people

“Cut-Fail” Combination II

- Let us define a predicate for the average tax rate
- However, there is a special tax rate for non-residents, i.e., they never pay the average rate

```
average_tax_rate(X) :- non_resident(X),fail.  
average_tax_rate(X) :- ...
```

- This will not work, as a non-resident will fail the first rule and then one of the following rules will be applied
- However, that's exactly what we don't want to happen
- The following will make sure that none of the following rules will be applied

```
average_tax_rate(X) :- non_resident(X),!,fail.  
average_tax_rate(X) :- ...
```

Generate and Test I

A common programming pattern in Prolog is “generate and test”

```
foo :- g1, g2, ..., gn, t1, t2, ..., tm.
```

The sequence of predicates $g1, g2, \dots, gn$ can succeed in many different ways

- They generate lots of different potential solutions

The sequence of predicates $t1, t2, \dots, tm$ tests whether something generated by $g1, g2, \dots, gn$ is actually a solution

- If something is not a solution, this causes $g1, g2, \dots, gn$ to backtrack and generate next candidate

Generate and Test II

Example: We want to define integer division just using addition and multiplication

- 1 Build a predicate that generates all integers:

```
is_integer(0).  
is_integer(X) :- is_integer(Y), X is Y+1.
```

- 2 Then we check the numbers generated by `is_integer`

```
idiv(X,Y,Result) :- is_integer(Result),  
                    Prod1 is Result*Y,  
                    Prod2 is (Result+1)*Y,  
                    Prod1 =< X, Prod2 > X,  
                    !.
```

Generate and Test III

- The first line in `idiv` is the generator, the other lines are implementing the test
- We know that there can only be one possible solution
- After reaching it, we can stop the search, otherwise `is_integer` would keep on producing potential `Results`

Cutting too Deeply I

- The cut operator is a dangerous tool and should be used sparingly
- It can behave in unexpected ways.
- We want to formulate that every person has two parents, except Adam and Eve who have no parents

```
parent(adam,0) :- !.  
parent(eve,0) :- !.  
parent(X,2).  
?- parent(eve,X).  
X = 0  
?- parent(john,X).  
X = 2  
?- parent(eve,2).  
yes
```

Cutting too Deeply II

- It is considered good programming style to replace cuts by the use of negation (if possible)

```
parent(adam,0).  
parent(eve,0).  
parent(X,2) :- \+(X = adam), \+(X = eve).  
?- parent(eve,X).  
X = 0 ? ;  
no  
?- parent(john,X).  
X = 2  
yes  
?- parent(eve,2).  
no
```

Cutting too Deeply III

The program computing the sum of the numbers from 1 to N can also be rewritten:

```
sum_to(N,1) :- N =< 1.  
sum_to(N,Result) :- N > 1,  
                    TmpN is N-1,  
                    sum_to(TmpN,TmpRes),  
                    Result is TmpRes + N.
```

This also makes it clear which rule to use when

Summary I

Strengths of Prolog

- Prolog is very well suited for application centered around Artificial Intelligence (AI)
 - Natural-language processing
 - AI behavior in games
 - Constraint satisfaction problems, such as time tabling and scheduling
- Prolog (or its descendants) is used in the context of the Semantic Web
 - A variant called Datalog is used in databases
- Also used for simulation and prediction software

Summary II

Weaknesses of Prolog

- Prolog has a steeper learning curve compared to other languages
- Fairly focused niche applications, not really a general-purpose language
- There are scalability issues, the basic matching strategy used by Prolog is computationally expensive
 - Has problems to process large data sets
- It is not as declarative as it seems at first glance
 - If you want to write efficient Prolog programs, you have to know what is going on behind the scenes

Questions...